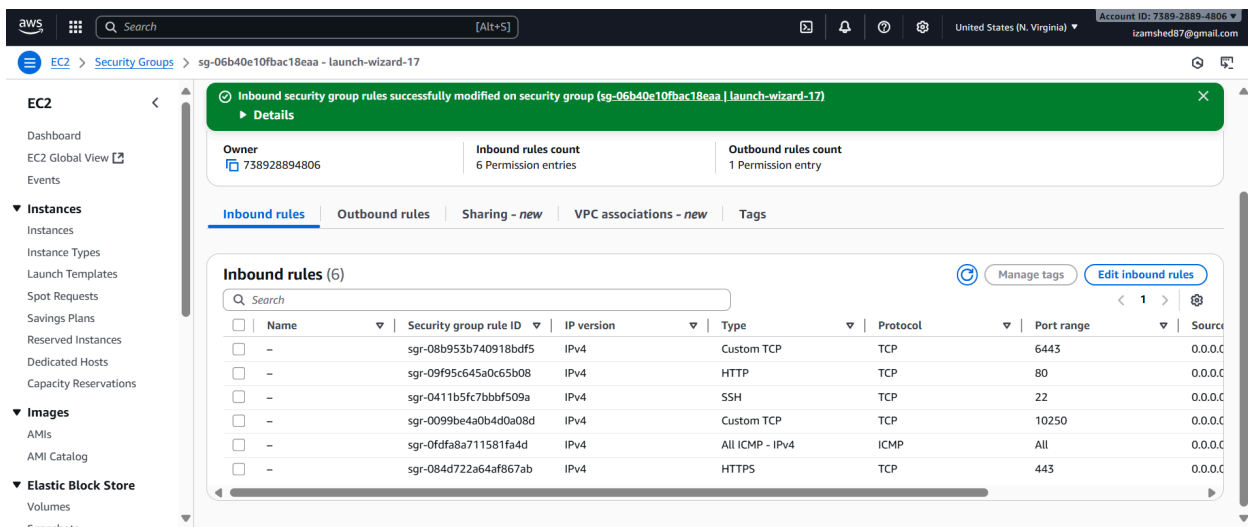


Before kubeadm join command allow these ports in the master node security group



1 Enable IP forwarding immediately

```
sudo sysctl -w net.ipv4.ip_forward=1
```

2 Make it persistent across reboots

Edit `/etc/sysctl.conf` (or create a drop-in file):

```
sudo nano /etc/sysctl.conf
```

Add (or uncomment) this line:

```
net.ipv4.ip_forward = 1
```

Then reload:

```
sudo sysctl -p
```

3 Verify

```
cat /proc/sys/net/ipv4/ip_forward  
# It should print: 1
```

4 Retry kubeadm

```
sudo kubeadm init \  
  --apiserver-advertise-address 172.31.28.30 \  
  --pod-network-cidr 10.244.0.0/16 \  
  --upload-certs
```

Tokens expire after 24h by default; regenerate if needed:

```
sudo kubeadm token create --print-join-command
```

1. Copy the master kubeconfig to the worker:

On the master:

```
sudo cat /etc/kubernetes/admin.conf
```

Set up kubeconfig:

```
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Copy its content to the worker and save as `$HOME/admin.conf` on the worker.

2. Set ownership and environment variable on the worker:

effect — you already have the **master's admin.conf** copied to the worker as `$HOME/admin.conf`.

Now you just need to **point kubectl on the worker to that file**:

```
export KUBECONFIG=$HOME/admin.conf  
echo 'export KUBECONFIG=$HOME/admin.conf' >> ~/.bashrc
```

After that, test the connection:

```
kubectl get nodes  
kubectl get pods -A
```

✅ You should see the **master node listed**.

```
kubeadm join 172.31.28.30:6443 --token fjayi9.s56fl9bpqyk2fpar --discovery-token-ca-cert-hash  
sha256:731fcbfc2a85c316724d92931514f9fff2473c6a400051a907703b28ebdaefdc
```

```
sudo kubeadm join 172.31.28.30:6443 \
```

```
--token fjayi9.s56fl9bpqyk2fpar \  
--discovery-token-ca-cert-hash  
sha256:731fcbfc2a85c316724d92931514f9fff2473c6a400051a907703b28ebdaefdc
```

```
ubuntu@Zamshed-Master-Node:~$ kubectl get no  
NAME                STATUS    ROLES    AGE      VERSION  
zamshed-master-node NotReady  control-plane  5h17m   v1.33.5  
zamshed-worker-node-1 NotReady  <none>      3h44m   v1.33.5  
ubuntu@Zamshed-Master-Node:~$ sudo kubeadm token create --print-join-command  
kubeadm join 172.31.28.30:6443 --token fjayi9.s56fl9bpqyk2fpar --discovery-token-ca-cert-hash sha256:  
731fcbfc2a85c316724d92931514f9fff2473c6a400051a907703b28ebdaefdc  
ubuntu@Zamshed-Master-Node:~$
```

```
ubuntu@Zamshed-Worker-Node-2:~$ kubectl get no  
NAME                STATUS    ROLES    AGE      VERSION  
zamshed-master-node NotReady  control-plane  5h21m   v1.33.5  
zamshed-worker-node-1 NotReady  <none>      3h48m   v1.33.5  
zamshed-worker-node-2 NotReady  <none>      67s     v1.33.5  
ubuntu@Zamshed-Worker-Node-2:~$
```

Here's the **completed, polished documentation** for that section.

It walks through all the prerequisites, network settings, `kubeadm init`, copying `kubeconfig`, and finally joining the worker.

Kubernetes Cluster Setup – Master & Worker Node Join

1 Prepare the Master Node

a) Open Required Ports in Security Group

In your cloud provider (e.g., AWS), edit the **Security Group** for the master node and allow inbound:

Port	Protocol	Purpose
6443	TCP	Kubernetes API server
2379–2380	TCP	etcd server client API
10250	TCP	Kubelet API
10251	TCP	kube-scheduler
10252	TCP	kube-controller-manager
10255 (optional)	TCP	Read-only Kubelet API

Also allow ICMP (ping) for troubleshooting.

b) Enable IP Forwarding

Kubernetes requires IP forwarding to be enabled.

Enable immediately

```
sudo sysctl -w net.ipv4.ip_forward=1
```

Make it persistent

Edit `/etc/sysctl.conf`:

```
sudo nano /etc/sysctl.conf
```

Add or uncomment:

```
net.ipv4.ip_forward = 1
```

Reload settings:

```
sudo sysctl -p
```

Verify

```
cat /proc/sys/net/ipv4/ip_forward
```

Expected output: 1

2 Initialize the Control Plane

On the master node, run:

```
sudo kubeadm init \
  --apiserver-advertise-address=172.31.28.30 \
  --pod-network-cidr=10.244.0.0/16 \
  --upload-certs
```

Save the `kubeadm join` command from the output. Tokens expire after 24h — if needed, regenerate:

```
sudo kubeadm token create --print-join-command
```

3 Configure `kubect1` on the Master

Set up `kubeconfig` so you can run `kubect1` as a normal user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Test:

```
kubect1 get nodes
```

4 Copy Kubeconfig to Worker (Optional for `kubect1`)

If you want to run `kubect1` from the worker:

On the master:

```
sudo cat /etc/kubernetes/admin.conf
```

1. Copy the content.

On the worker:

```
nano ~/admin.conf    # paste the content
export KUBECONFIG=$HOME/admin.conf
echo 'export KUBECONFIG=$HOME/admin.conf' >> ~/.bashrc
Verify:
```

```
kubectl get nodes
kubectl get pods -A
```

- 2.

✅ You should see the master node.

5 Join the Worker Node

Use the join command saved from the `kubeadm init` step.

Run it with **sudo**:

```
sudo kubeadm join 172.31.28.30:6443 \
  --token fjayi9.s56fl9bpqyk2fpar \
  --discovery-token-ca-cert-hash
sha256:731fcbfc2a85c316724d92931514f9fff2473c6a400051a907703b28ebdaefd
c
```

⚠️ If you forget `sudo`, you'll see:

```
[ERROR IsPrivilegedUser]: user is not running as root
```

6 Verify the Node Status

On the master node:

```
kubectl get nodes -o wide
```

Example output:

NAME	STATUS	ROLES	AGE	VERSION
zamshed-master-node	Ready	control-plane	10m	v1.33.5
zamshed-worker-node-1	NotReady	<none>	1m	v1.33.5

Nodes may show **NotReady** until you deploy a CNI plugin.

7 Deploy a Pod Network (CNI)

Install a network plugin, e.g., Flannel:

```
kubectl apply -f  
https://raw.githubusercontent.com/flannel-io/flannel/master/Documentation/kube-flannel.yml
```

Check:

```
kubectl get pods -A  
kubectl get nodes
```

Once the CNI is ready, worker nodes should move to **Ready**.

Final Checklist

- Security group allows required ports
- IP forwarding enabled
- `kubeadm init` completed

- `kubeconfig` set up
- Worker joined successfully
- CNI deployed

Here's a single `sed` command to do that:

```
sed -i '/- --secure-port=10250/c\          - --secure-port=4443\n- --kubelet-i
```

```
ubuntu@controlpanel: ~$ hostname -I
172.31.23.139
ubuntu@controlpanel:~$ sudo kubeadm init \
--apiserver-advertise-address 172.31.23.139 \
--pod-network-cidr "10.244.0.0/16" \
--upload-certs
[init] Using Kubernetes version: v1.34.1
[preflight] Running pre-flight checks
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action beforehand using 'kubeadm config images pull'
W0916 17:28:36.799849    2236 checks.go:830] detected that the sandbox image "registry.k8s.io/pause:3.8" of the container ru
ntime is inconsistent with that used by kubeadm. It is recommended to use "registry.k8s.io/pause:3.10.1" as the CRI sandbox
image.

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.23.139:6443 --token ua67xd.qu210oltaakskud6 \
--discovery-token-ca-cert-hash sha256:6df2be6eaf7dbcbac545329c33a3d07dc6dc5d8f866500408a890ceb5c38aa56
ubuntu@controlpanel:~$ mkdir -p $HOME/.kube
ubuntu@controlpanel:~$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
ubuntu@controlpanel:~$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
ubuntu@controlpanel:~$
```



```
kubectl -n kube-system rollout status deploy metrics-server  
kubectl get apiservices | grep metrics  
kubectl top nodes  
kubectl top pods
```